

SePRo: Separation Process with R

Estimating of parameters in statistical distributions using R. Package provides an implementation of the EM algorithm (Expectation Maximization Algorithm) in the R language. Today, EM and its variants are regularly used to solve a broad range of today's estimation problems, from the multiple EM for motif elicitation (MEME) algorithm for motif-finding in DNA sequences, to fitting mixture models to disambiguate targets from clutter in radar. Hope that you, too, will find EM useful. You can see details of usage [here](#).




Figure 1: *Example of mixture - 2 Gaussian functions*

Installation in R

sepro is a GitHub package so you can use 'install_github()' from **devtools** package. Install **devtools** first:

```
1 if("devtools" %in% rownames(installed.packages())) == FALSE){  
2   install.packages("devtools")  
3 }  
4 library(devtools)
```

After that you can install **sepro** package:

```
1 install_github("hdrbv/sepro", ref = "main")  
2 library(sepro)
```

Theory

Let's assume that we have some observed data y , a parametric density $p(y|\theta)$, a description of some complete data x that we wish we had, and the parametric density $p(x|\theta)$. We assume that the complete data can be modeled as a continuous random vector X with density $p(x|\theta)$, where $\theta \in \Omega$ for some set Ω . We do not observe X directly, instead we observe a realization y of the random vector Y that depends on X . For example, Y might be the first component of the vector X .

Given that we only have y , the main goal here is to find the maximum likelihood estimate (MLE) of θ :

$$\hat{\theta}_{MLE} = \operatorname{argmax} p(y|\theta)$$

It's often easier to calculate the θ that maximizes the log-likelihood of y :

$$\hat{\theta}_{MLE} = \operatorname{argmax} \log p(y|\theta)$$

Because $\log()$ is a monotonically increasing function, solutions will be the same for both equations. But sometimes it's difficult to solve them. Then we can try EM: we make a guess about the complete data X and solve for the θ that maximizes the (expected) log-likelihood of X . And once we have an estimate for θ , we can make a better guess about the complete data X , and iterate.

Let's break E-step and M-step of algorithm down into five steps:

1. Let $m = 0$ and make initial estimate $\theta^{(m)}$ for θ
2. Given the observed data y and pretending for the moment that your current guess $\theta^{(m)}$ is correct, formulate the conditional probability distribution $p(x|y, \theta^{(m)})$ for the complete data x
3. Using the conditional probability distribution $p(x|y, \theta^{(m)})$ calculated in step 2, form the conditional expected log-likelihood, which is called the Q-function:

$$\begin{aligned} Q(\theta|\theta^m) &= \int \log p(x|\theta) p(x|y, \theta^m) dx = \\ &= E_{X|y, \theta^m}(\log p(X|\theta)) \end{aligned}$$

where the integral is over set $\chi(y)$, which is the closure of the set $\{x|p(x|y, \theta) > 0\}$, and assume that $\chi(y)$ does not depend on θ .

4. Find the θ that maximizes Q - function; result = our new estimate = $\theta^{(m+1)}$
5. Let $m = m + 1$ and go back to Step №2. EM algorithm does not specify a stopping criterion; standard criteria are to iterate until the estimate stops changing: $|\theta^{(m+1)} - \theta^{(m)}| < \epsilon$ for some $\epsilon > 0$, or to iterate until the log-likelihood $l(\theta^{m+1}) - l(\theta^m) < \epsilon$ for some $\epsilon > 0$

EM algorithm is only guaranteed to never get worse. Usually, it will find a peak in the likelihood $p(y|\theta)$, but if the likelihood function $p(y|\theta)$ has multiple peaks, EM will not necessarily find the global maximum of the likelihood. In practise, it's common to start EM from multiple random initial guesses, and choose the one with the largest likelihood as the final guess for θ

Practice

Let's apply theory to practice and also check the work of [sepro](#) package.

Import additional packages:

```
1 if("pacman" %in% rownames(installed.packages()) == FALSE){  
2   install.packages("pacman")}  
3 pacman:: p_load(dplyr, ggplot2, stargazer)
```

Firstly, let's create mixture of two distributions which we will separate:

```
1 set.seed(1) #fix results of randomization  
2 cond <- sample(c(0, 1), size = 500,  
3               replace = TRUE, prob = c(0.4, 0.6))  
4 # Sample from two different Gaussian distributions  
5 mix <- ifelse(cond == 1, rnorm(n = 500, mean = 5, sd = 1.5),  
6               rnorm(n = 500, mean = 0, sd = 1))  
7 plot(mix)
```




Figure 2: *Plot of created mixture (result of `plot(mix)`)*

Apply *EM* function from **sepro** package to our mixture:

```
1 vect <- as.numeric(mix)
2 EM1 <- EM(vect , 2)
```

And use *plot_em* function from package to see results of separation process:

```
1 plot_em(vect , EM1)
```




Figure 3: *Plot of separated mixture (result of plot_em)*

That's it. We have a quite accurate parameter estimation of our distributions - it's really close to real:

Table 1

	Expected value	Variance
Distribution 1 (from initial dataset)	0	1
Distribution 2 (from initial dataset)	5	1.5
Distribution 1 (after separation process)	0	1
Distribution 2 (after separation process)	5	2